

Chapter 9: Random Forest Classifier

(c) 2019 Galit Shmueli, Peter C. Bruce, Peter Gedeck

Code included in

Data Mining for Business Analytics: Concepts, Techniques, and Applications in Python (First Edition) Galit Shmueli, Peter C. Bruce, Peter Gedeck, and Nitin R. Patel. 2019.

Import required packages

In [28]:

```
%matplotlib inline

from pathlib import Path

import pandas as pd
from sklearn import preprocessing
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn import metrics
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
import matplotlib.pyplot as plt
from dmba import classificationSummary, gainsChart, liftChart
```

This is the same dataset as used in Homework #3.

In [29]:

```
hr_df = pd.read_csv('/Users/aminazimi/Downloads/WA_Fn-UseC_-HR-Employee-Attrition.csv')
hr_df.head()
```

Out[29]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeM
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeN
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	

5 rows × 35 columns

In [30]: `hr_df.shape`

Out[30]: (1470, 35)

In [31]: `hr_df.dtypes`

Out[31]:

Age	int64
Attrition	object
BusinessTravel	object
DailyRate	int64
Department	object
DistanceFromHome	int64
Education	int64
EducationField	object
EmployeeCount	int64
EmployeeNumber	int64
EnvironmentSatisfaction	int64
Gender	object
HourlyRate	int64
JobInvolvement	int64
JobLevel	int64
JobRole	object
JobSatisfaction	int64
MaritalStatus	object
MonthlyIncome	int64
MonthlyRate	int64
NumCompaniesWorked	int64
Over18	object
OverTime	object

```

PercentSalaryHike      int64
PerformanceRating      int64
RelationshipSatisfaction int64
StandardHours          int64
StockOptionLevel       int64
TotalWorkingYears      int64
TrainingTimesLastYear  int64
WorkLifeBalance        int64
YearsAtCompany         int64
YearsInCurrentRole     int64
YearsSinceLastPromotion int64
YearsWithCurrManager   int64
dtype: object

```

```

In [32]: # Create a y response variable and an X collection of predictors
y = hr_df['Attrition']
X = hr_df.drop(columns=['Attrition'])
print(len(X.columns))

```

34

```

In [33]: # Dummy code in preparation of logistic regression
X = pd.get_dummies(X, prefix_sep='_', drop_first=True)
print(len(X.columns))

```

47

```

In [34]: # Convert the text of Gone to a binary numeric variable (0/1)
y = y.astype('category').cat.codes

# Check for a class imbalance
y.value_counts()

```

```

Out[34]: 0    1233
         1     237
         dtype: int64

```

Print out a list of attributes by name and sequence number to prepare for ADASYN

```

In [35]: print(pd.DataFrame(X.columns))

```

0

0
Age

1	DailyRate
2	DistanceFromHome
3	Education
4	EmployeeCount
5	EmployeeNumber
6	EnvironmentSatisfaction
7	HourlyRate
8	JobInvolvement
9	JobLevel
10	JobSatisfaction
11	MonthlyIncome
12	MonthlyRate
13	NumCompaniesWorked
14	PercentSalaryHike
15	PerformanceRating
16	RelationshipSatisfaction
17	StandardHours
18	StockOptionLevel
19	TotalWorkingYears
20	TrainingTimesLastYear
21	WorkLifeBalance
22	YearsAtCompany
23	YearsInCurrentRole
24	YearsSinceLastPromotion
25	YearsWithCurrManager
26	BusinessTravel_Travel_Frequently
27	BusinessTravel_Travel_Rarely
28	Department_Research & Development
29	Department_Sales
30	EducationField_Life Sciences
31	EducationField_Marketing
32	EducationField_Medical
33	EducationField_Other
34	EducationField_Technical Degree
35	Gender_Male
36	JobRole_Human Resources
37	JobRole_Laboratory Technician
38	JobRole_Manager
39	JobRole_Manufacturing Director
40	JobRole_Research Director
41	JobRole_Research Scientist
42	JobRole_Sales Executive
43	JobRole_Sales Representative
44	MaritalStatus_Married

```
45         MaritalStatus_Single
46         OverTime_Yes
```

Train/test split with stratification of the response variable

```
In [36]: # Split the data into training and test sets (holdout approach)
train_X, test_X, train_y, test_y = train_test_split(X, y, test_size=0.2, stratify=y, random_state=2019)
```

Fix the class imbalance issue on the training data.

```
In [37]: from imblearn.over_sampling import ADASYN

ada = ADASYN()
train_X, train_y = ada.fit_sample(train_X, train_y.ravel())
```

Logistic Regression Baseline

```
In [38]: # Build a logistic regression model as a baseline
logit_reg = LogisticRegression(penalty="l2", C=1e42, solver='liblinear')
logit_reg.fit(train_X, train_y)
```

```
Out[38]: LogisticRegression(C=1e+42, solver='liblinear')
```

```
In [39]: # we are only interested in classification accuracy
classificationSummary(train_y, logit_reg.predict(train_X))
classificationSummary(test_y, logit_reg.predict(test_X))
```

Confusion Matrix (Accuracy 0.9142)

	Prediction	
Actual	0	1
0	930	56
1	110	839

Confusion Matrix (Accuracy 0.8367)

	Prediction	
Actual	0	1
0	226	21
1	27	20

```
In [40]: classes = logit_reg.predict(test_X)

print(metrics.classification_report(test_y, classes))
```

	precision	recall	f1-score	support
0	0.89	0.91	0.90	247
1	0.49	0.43	0.45	47
accuracy			0.84	294
macro avg	0.69	0.67	0.68	294
weighted avg	0.83	0.84	0.83	294

Use RandomForest

Start by recreating the X and y objects to change one-hot encoding parameter

```
In [41]: # Create a y response variable and an X collection of predictors
y = hr_df['Attrition']
X = hr_df.drop(columns=['Attrition'])
```

```
In [42]: # Dummy code in preparation of RandomForest model
X = pd.get_dummies(X, prefix_sep='_', drop_first=False)
print(len(X.columns))
```

55

```
In [43]: # Convert the text of Gone to a binary numeric variable (0/1)
y = y.astype('category').cat.codes
y.value_counts()
```

```
Out[43]: 0    1233
1      237
dtype: int64
```

```
In [44]: # Split the data into training and test sets (holdout approach)
train_X, test_X, train_y, test_y = train_test_split(X, y, test_size=0.2, stratify=y, random_state=1)
```

```
In [45]: from imblearn.over_sampling import ADASYN
```

```
ada = ADASYN()
train_X, train_y = ada.fit_sample(train_X, train_y)
```

In [46]:

```
# user grid search to find optimized tree
param_grid = {
    'n_estimators': [550, 600, 650],
    'criterion' : ['entropy', 'gini'],
    'oob_score': ['True'],
    'min_impurity_decrease': [0.0001, .0005, 0.001],
    'min_samples_split': [2, 4, 6, 8, 10, 12, 14],
}
```

In [47]:

```
gridSearch = GridSearchCV(RandomForestClassifier(), param_grid, cv=2, n_jobs=-1)

gridSearch.fit(train_X, train_y)

print('Initial parameters: ', gridSearch.best_params_)

rfTree = gridSearch.best_estimator_
```

Initial parameters: {'criterion': 'entropy', 'min_impurity_decrease': 0.0005, 'min_samples_split': 2, 'n_estimators': 650, 'oob_score': 'True'}

Based on these chosen hyperparameters, reprogram the GridSearchCV for a finer search pattern and run again

In [48]:

```
print(rfTree.oob_score_)
```

0.9363683393688567

In [49]:

```
# we are only interested in classification accuracy
classificationSummary(train_y, rfTree.predict(train_X))

classificationSummary(test_y, rfTree.predict(test_X))
```

Confusion Matrix (Accuracy 1.0000)

	Prediction	
Actual	0	1
0	986	0
1	0	947

Confusion Matrix (Accuracy 0.8639)

	Prediction	
Actual	0	1
0	236	11
1	29	18

In [50]:

```
classes = rfTree.predict(test_X)

print(metrics.classification_report(test_y, classes))
```

	precision	recall	f1-score	support
0	0.89	0.96	0.92	247
1	0.62	0.38	0.47	47
accuracy			0.86	294
macro avg	0.76	0.67	0.70	294
weighted avg	0.85	0.86	0.85	294

In [51]:

```
%matplotlib inline
import numpy as np

train_X = pd.DataFrame(train_X)

importances = rfTree.feature_importances_
std = np.std([tree.feature_importances_ for tree in rfTree.estimators_], axis=0)

df = pd.DataFrame({'feature': train_X.columns, 'importance': importances, 'std': std})
df = df.sort_values('importance')
print(df)

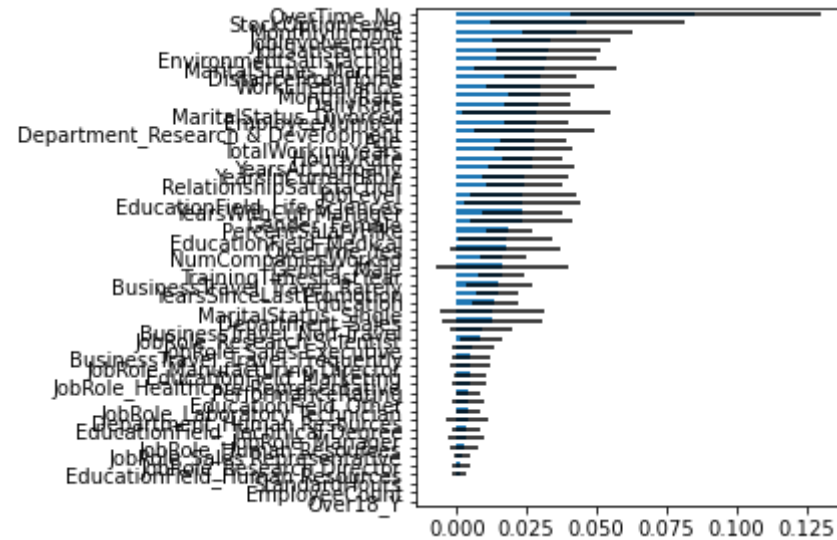
ax = df.plot(kind='barh', xerr='std', x='feature', legend=False)
ax.set_ylabel('')

plt.tight_layout()
plt.show()
```

	feature	importance	std
52	Over18_Y	0.000000	0.000000
4	EmployeeCount	0.000000	0.000000
17	StandardHours	0.000000	0.000000
32	EducationField_Human Resources	0.001050	0.002481

45	JobRole_Research Director	0.001477	0.003329
48	JobRole_Sales Representative	0.001953	0.003030
41	JobRole_Human Resources	0.002479	0.005130
43	JobRole_Manager	0.003424	0.006331
37	EducationField_Technical Degree	0.003461	0.004841
29	Department_Human Resources	0.003739	0.007285
42	JobRole_Laboratory Technician	0.003799	0.004254
36	EducationField_Other	0.003828	0.005746
15	PerformanceRating	0.004056	0.004391
40	JobRole_Healthcare Representative	0.004466	0.006217
34	EducationField_Marketing	0.004548	0.005844
44	JobRole_Manufacturing Director	0.004581	0.007037
27	BusinessTravel_Travel_Frequently	0.005003	0.006703
47	JobRole_Sales Executive	0.005813	0.007366
46	JobRole_Research Scientist	0.008602	0.007331
26	BusinessTravel_Non-Travel	0.008962	0.011186
31	Department_Sales	0.012481	0.017998
51	MaritalStatus_Single	0.012788	0.018689
3	Education	0.013494	0.008160
24	YearsSinceLastPromotion	0.014650	0.007634
28	BusinessTravel_Travel_Rarely	0.015051	0.011942
20	TrainingTimesLastYear	0.015918	0.008534
39	Gender_Male	0.016353	0.023789
13	NumCompaniesWorked	0.016583	0.008010
54	OverTime_Yes	0.017501	0.019800
35	EducationField_Medical	0.017555	0.016783
14	PercentSalaryHike	0.018481	0.008281
38	Gender_Female	0.023157	0.018143
25	YearsWithCurrManager	0.023293	0.014175
33	EducationField_Life Sciences	0.023576	0.020836
9	JobLevel	0.023641	0.018790
16	RelationshipSatisfaction	0.023903	0.013635
23	YearsInCurrentRole	0.024322	0.015407
22	YearsAtCompany	0.026923	0.015451
7	HourlyRate	0.027056	0.010976
19	TotalWorkingYears	0.027402	0.013848
0	Age	0.027570	0.011730
30	Department_Research & Development	0.027627	0.021724
5	EmployeeNumber	0.028345	0.011215
49	MaritalStatus_Divorced	0.028383	0.026416
1	DailyRate	0.029055	0.011895
12	MonthlyRate	0.029524	0.011369
21	WorkLifeBalance	0.029726	0.019256
2	DistanceFromHome	0.029826	0.012640
50	MaritalStatus_Married	0.031528	0.025325

6	EnvironmentSatisfaction	0.031860	0.017799
10	JobSatisfaction	0.032577	0.018582
8	JobInvolvement	0.033598	0.021211
11	MonthlyIncome	0.043121	0.019728
18	StockOptionLevel	0.046621	0.034801
53	OverTime_No	0.085268	0.044898



In [52]:

```

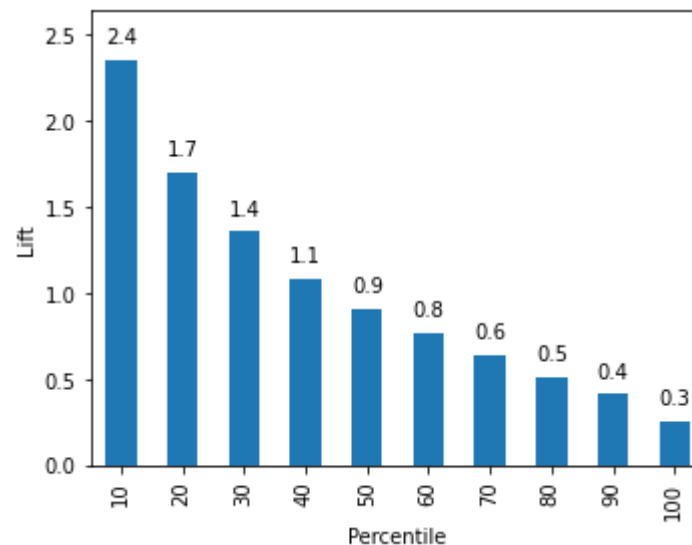
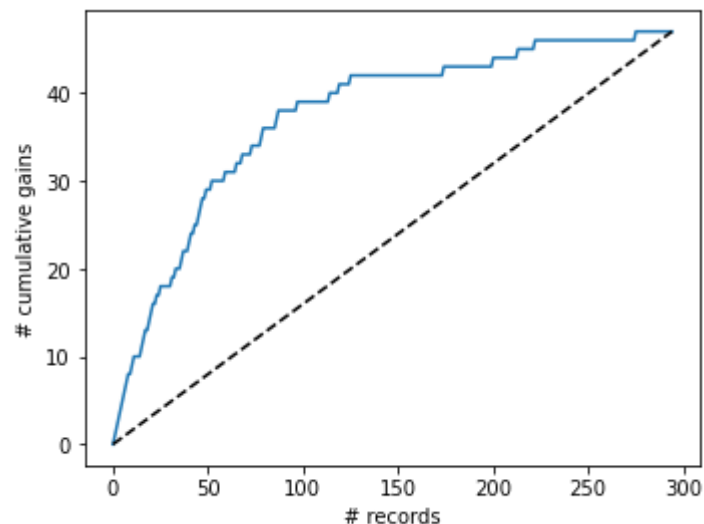
rfTree_pred = rfTree.predict(test_X)
rfTree_proba = rfTree.predict_proba(test_X)
rfTree_result = pd.DataFrame({'actual': test_y,
                              'p(0)': [p[0] for p in rfTree_proba],
                              'p(1)': [p[1] for p in rfTree_proba],
                              'predicted': rfTree_pred })

df = rfTree_result.sort_values(by=['p(1)'], ascending=False)
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(10, 4))

gainsChart(df.actual, ax=axes[0])
liftChart(df['p(1)'], title=False, ax=axes[1])

plt.tight_layout()
plt.show()

```



In [53]:

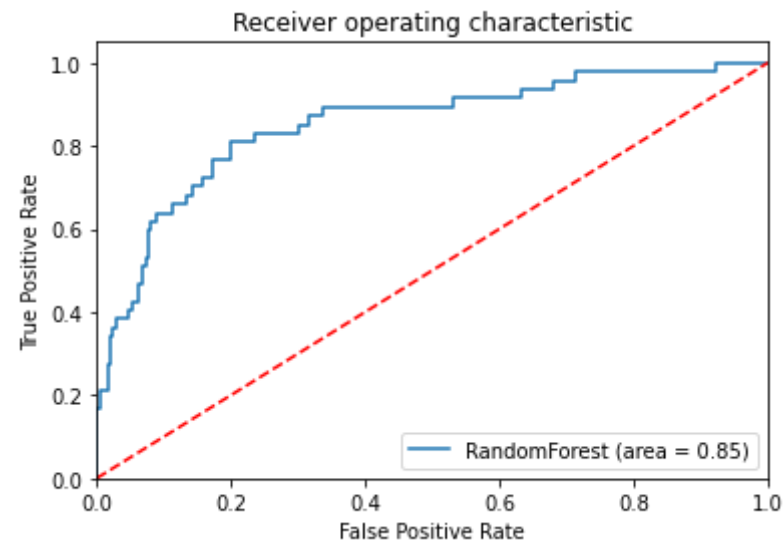
```

rfTree_pred = rfTree.predict(test_X)
rfTree_proba = rfTree.predict_proba(test_X)

preds = rfTree_proba[:,1]
fpr, tpr, threshold = metrics.roc_curve(test_y, preds)
roc_auc = metrics.auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, label='RandomForest (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('RFTree_ROC')
plt.show()

```



In []: